
Active Queue Managements In Residential Setting

Ish Kumar Jain
Supervisor: Prof. Shiv Panwar
and Fraida Fund

Goal: Reproducing the results of-



Contents lists available at [ScienceDirect](#)

Computer Networks

journal homepage: www.elsevier.com/locate/comnet



The Good, the Bad and the WiFi: Modern AQMs in a residential setting

Toke Høiland-Jørgensen*, Per Hurtig Anna Brunstrom

Department of Mathematics and Computer Science, Karlstad University, 651 88 Karlstad, Sweden



Problem:

**Linux Default:
FIFO Queuing**

BufferBloat

Network Bottleneck is congested.
Large buffer fill up and do not drain.




Image Source: <https://forum.peplink.com/t/bufferbloat/3331/21>

Solution: 3 AQMs

Adaptive RED

Adjust the RED max dropping probability based on observed queue length.

CoDel

Controlled Delay:
Directly measures the time pkt spent in a controlled queue.

PIE

Proportional Integral Controller: Infers Delay from instantaneous Queuing occupancy

But,

Can we do better than AQMs?

Fairness Queuing

SFQ: Hashing in sub-queues
Served as Round Robin

Fq_codel: FQ with CoDel

Fq_nocodel: FQ without AQM

Performance Metrics

Good

Steady State Behavior:
Throughput and Latency behavior.

Real-time Response under load (**RRUL Test**):
4 concurrent TCP with UDP and ICMP packets.

Bad

Transient Behavior:
AQM gives high latency spikes.

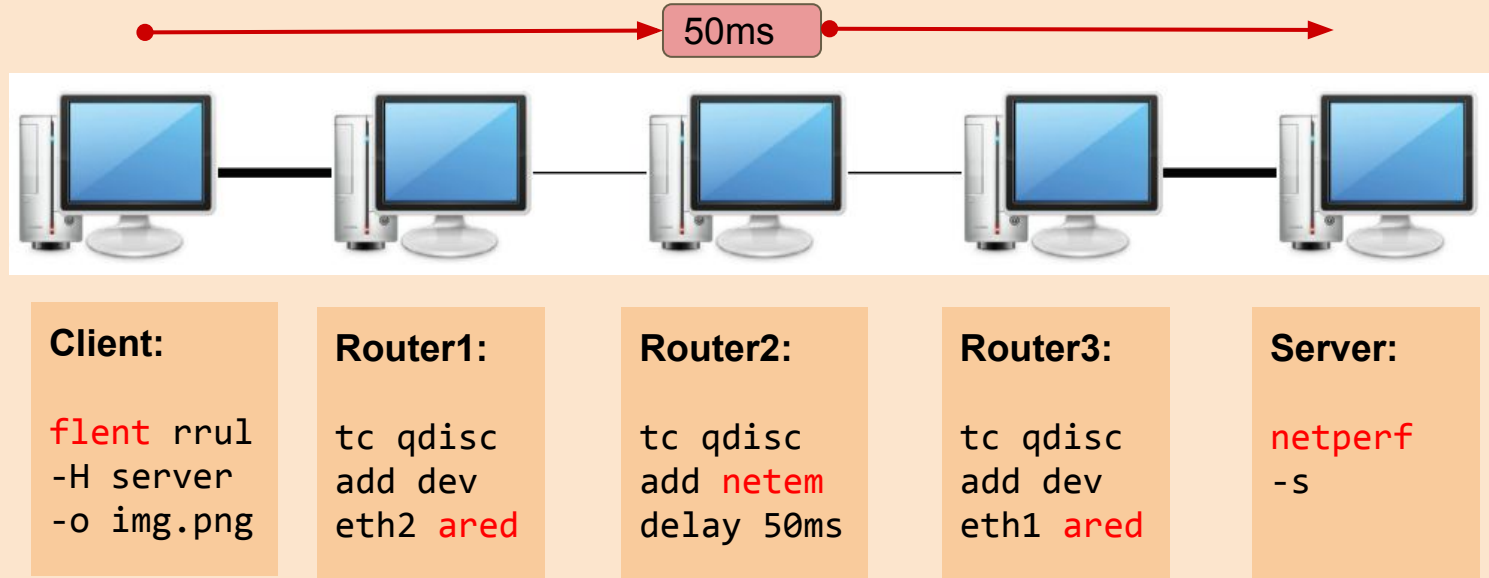
Fairness among flows:
AQM are not fairer

WiFi

Behavior on WiFi Link:

Performance drop as compared to the Wired link in all properties.

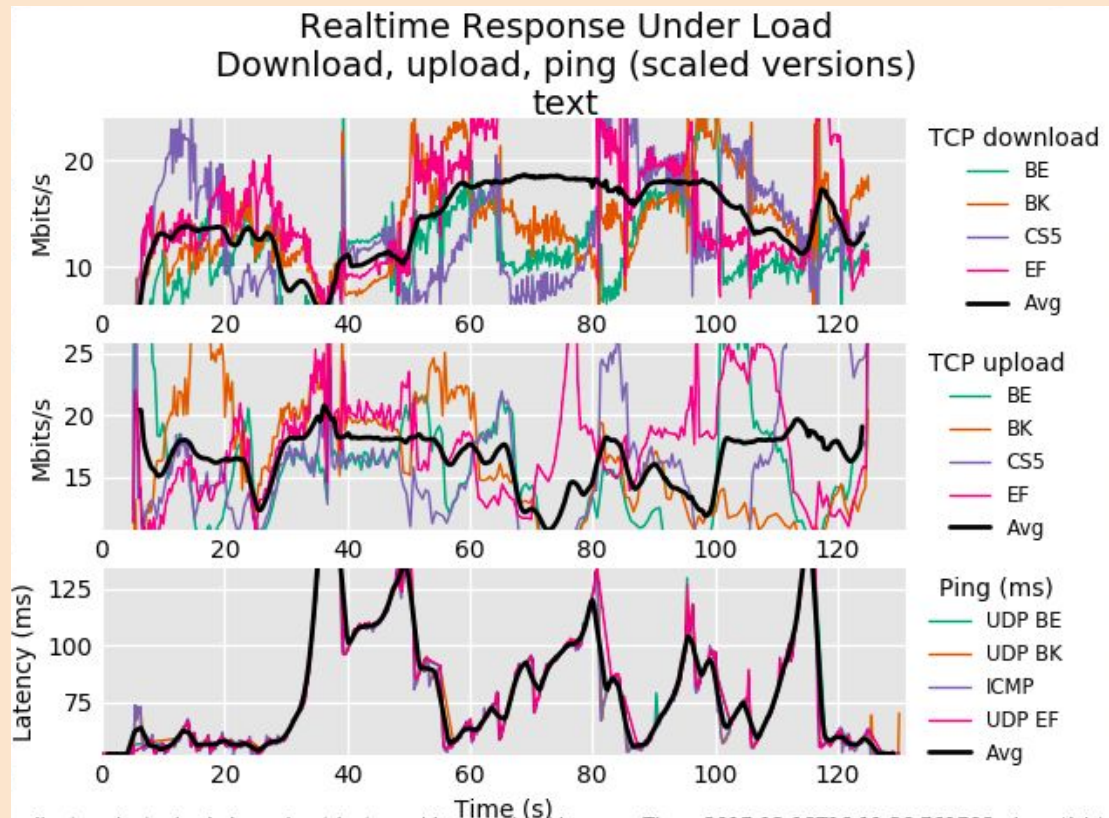
Implementation: GENI Testbed



Run my Experiment:

Install **flent** and **netperf** at client and server | Change **ssh** and **scp** commands in the code |
Run `bash maincode.sh 10` | Find data in a local directory | Analyze it on **Matlab**.

Flent: RRUL Test details



Four priority queues in WiFi

BE: Best Effort

BK: Background,

CS5:Class Selector 5 (video)

EF: Expedient Forwarding(voice)

Avg: Average of four flows.

Latency Measurement

Measure the average latency of four concurrent flows (UDP and ICMP)

AQM Parameters

- Many AQM parameters can be tuned.
- But, we use the default parameters for AQMs that have generally shown good results.
- Strict buffer size: 127 packets for most algorithms.
- Left: default parameters in *italic*

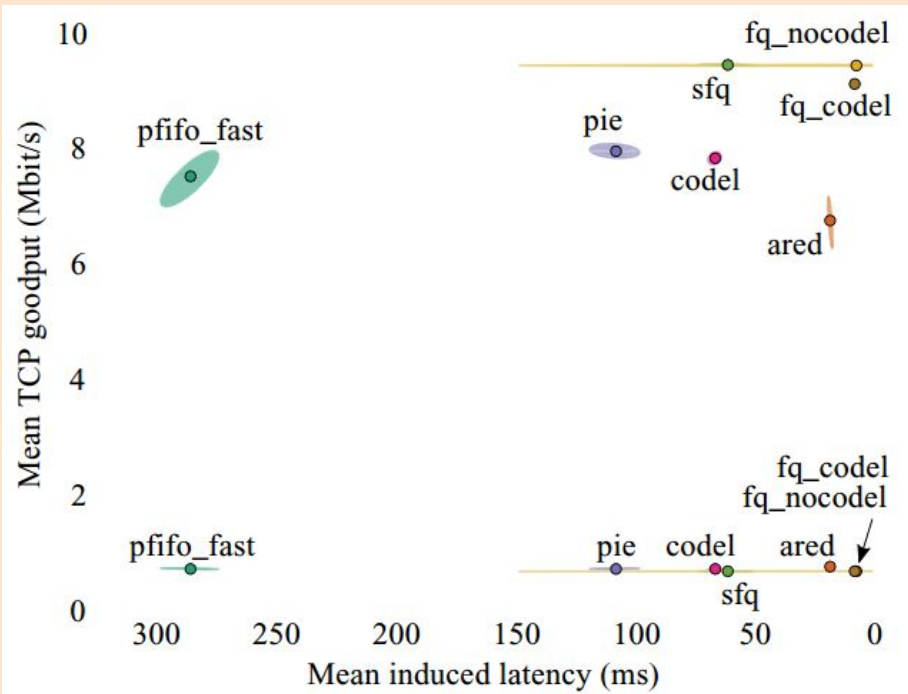
Parameter	1 Mbps	10 Mbps	100 Mbps
pfifo_fast			
txqueuelen	127	127	<i>1000</i>
ARED			
min	1514	12500	125000
bandwidth	1 Mbps	10 Mbps	100 Mbps
max	3028	–	–
PIE			
target	<i>20 ms</i>	<i>20 ms</i>	<i>20 ms</i>
tupdate	<i>30 ms</i>	<i>30 ms</i>	<i>30 ms</i>
limit	<i>1000</i>	<i>1000</i>	<i>1000</i>
CoDel			
target	<i>13 ms</i>	<i>5 ms</i>	<i>5 ms</i>
interval	<i>100 ms</i>	<i>100 ms</i>	<i>100 ms</i>
limit	<i>1000</i>	<i>1000</i>	<i>1000</i>
SFQ			
limit	<i>127</i>	<i>127</i>	<i>1000</i>
fq_codel			
target	<i>13 ms</i>	<i>5 ms</i>	<i>5 ms</i>
interval	<i>100 ms</i>	<i>100 ms</i>	<i>100 ms</i>
limit	<i>10240</i>	<i>10240</i>	<i>10240</i>
fq_nocodel			
limit	<i>127</i>	<i>127</i>	<i>1000</i>
interval	<i>100 s</i>	<i>100 s</i>	<i>100 s</i>

Implementation details

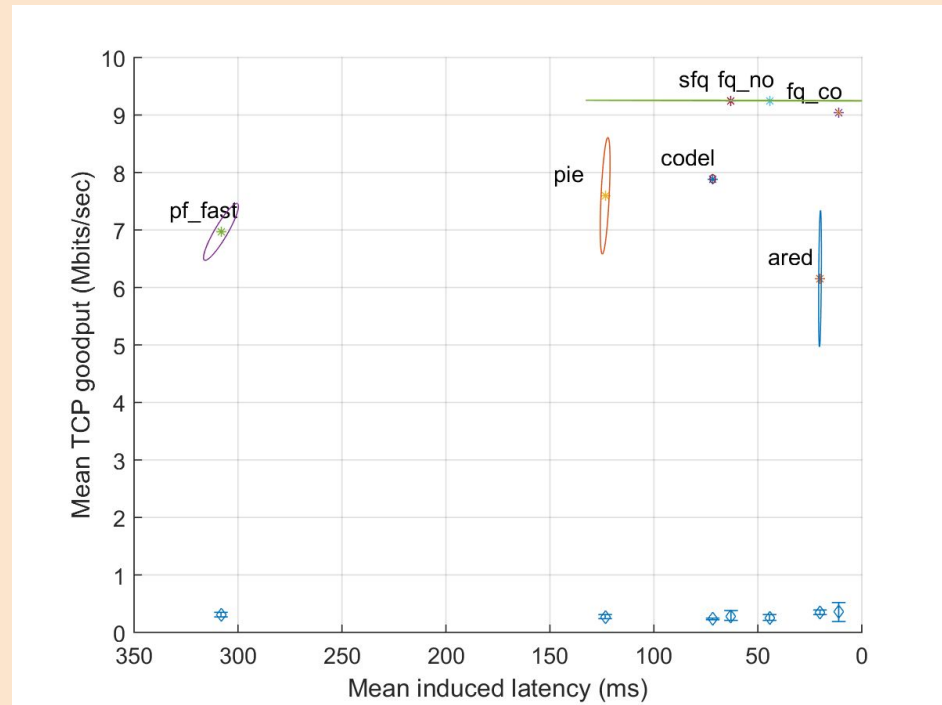
Required by Paper	My Implementation	Description
Use Dummynet to generate delay	Used netem instead	Dummynet not compatible
Run Flent and netperf for both RRUL and fairness test.	Flent used for RRUL test. But, iperf for fairness test	Due to difficulties in implementing multiple flows with Flent.
Repeat 30 times	Repeated 5 times	It takes a lot of time.
Hardware offload feature turn off	Turned off	On client: <code>ethtool -K interface tso off gso off</code>
Kernel clock frequency 1000 Hz	250 Hz	Caused issue at 100 Mbps
Tcp default cubic algorithm	Used cubic	Other TCP algos give same results
Fairness test on 4 concurrent flows	Implemented only 3 flows	Netem didn't allow >3 firewall rule

Results: RRUL Test 10/1

Our results are **comparable** to that of the paper. A **small difference in magnitude** can be attributed to some of the **implementation issues** discussed previously.

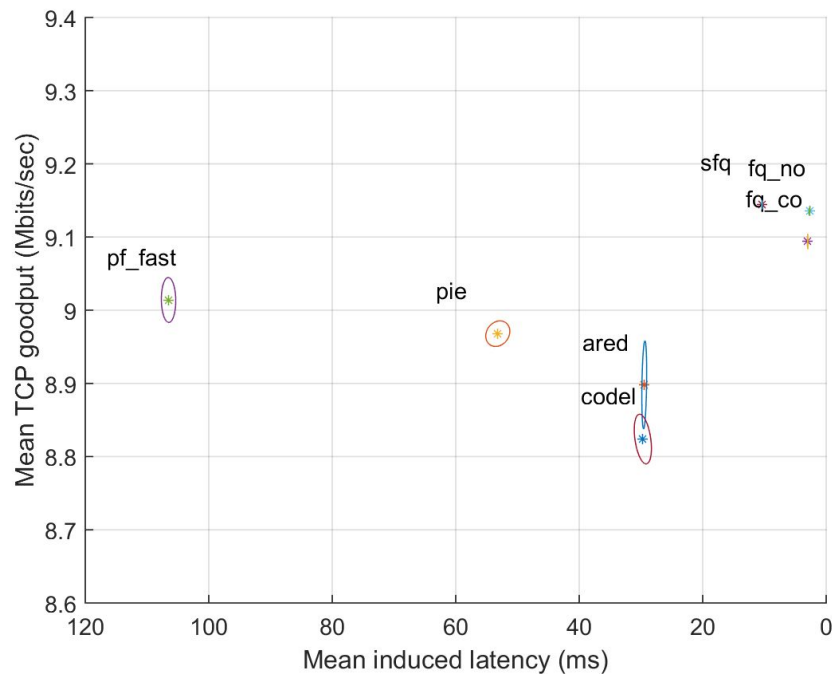
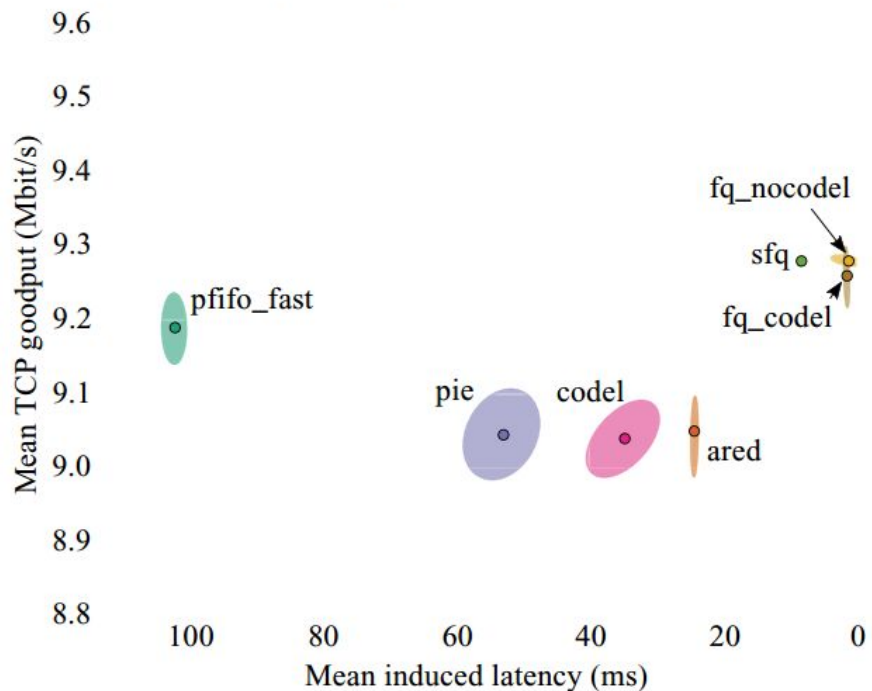


Toke et.al.

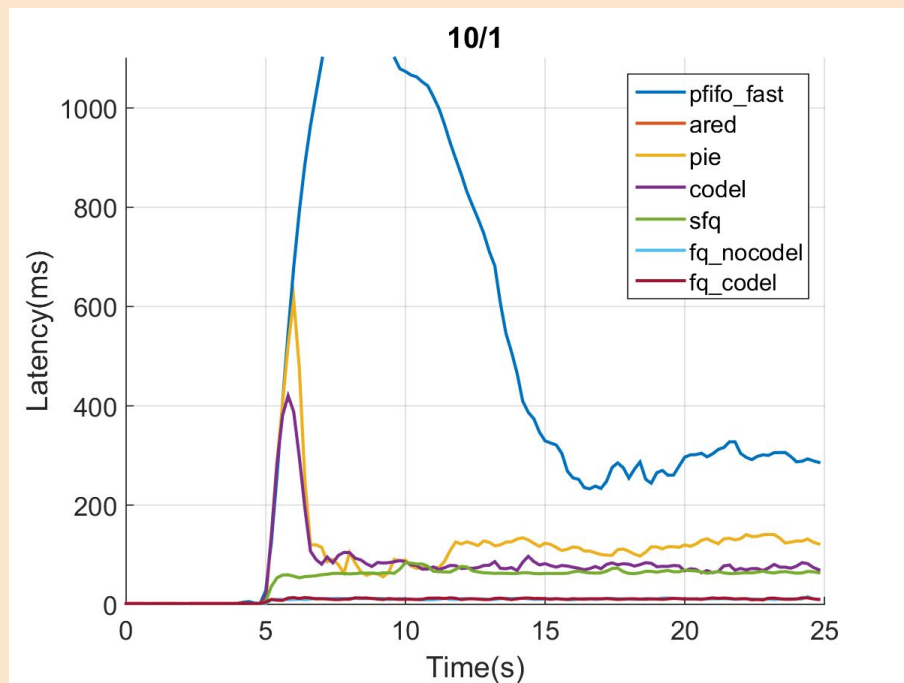
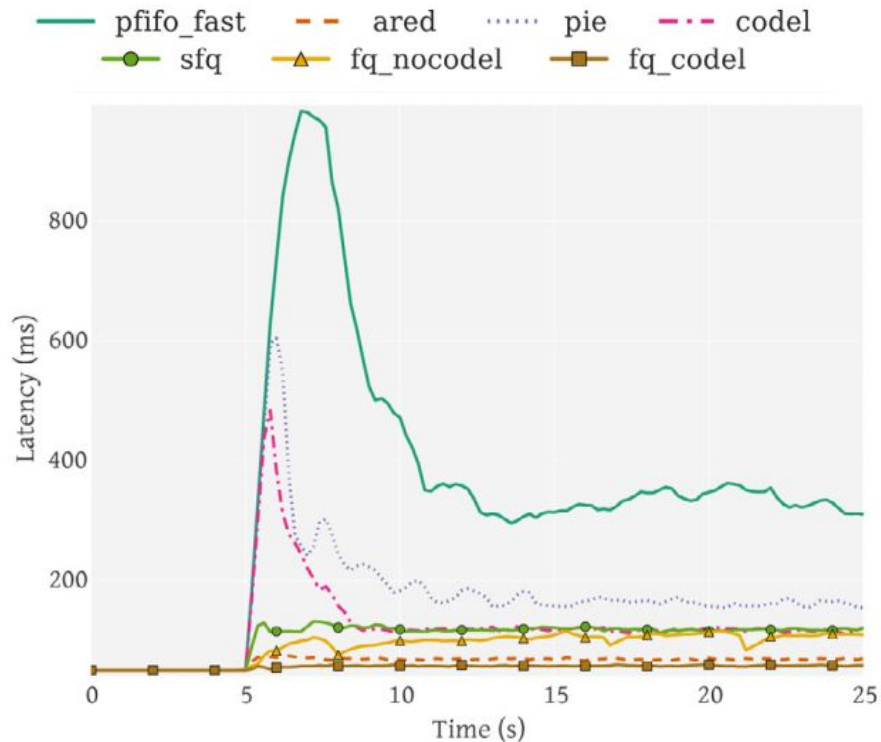


Our Results

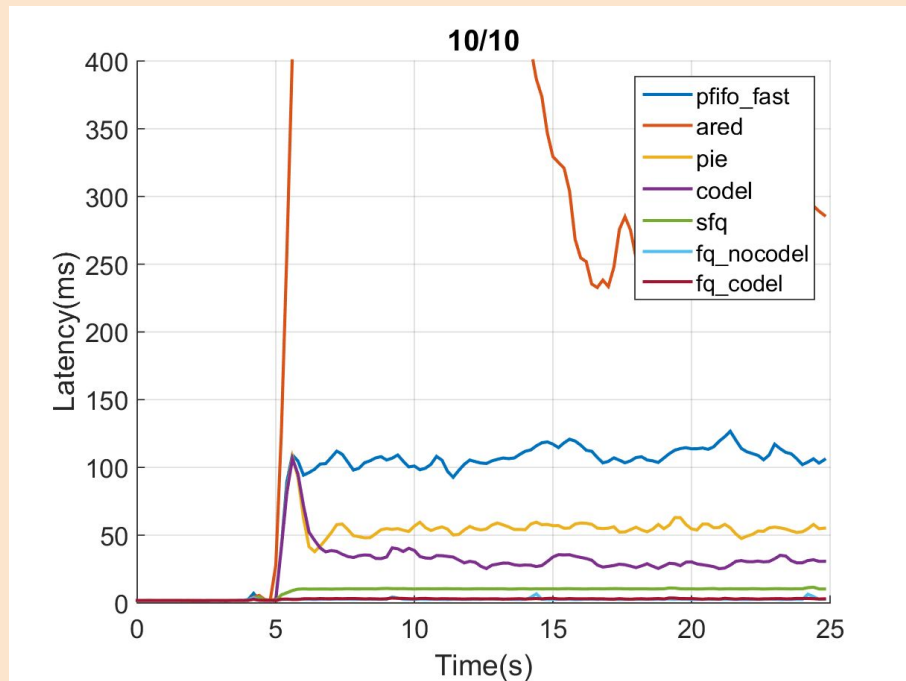
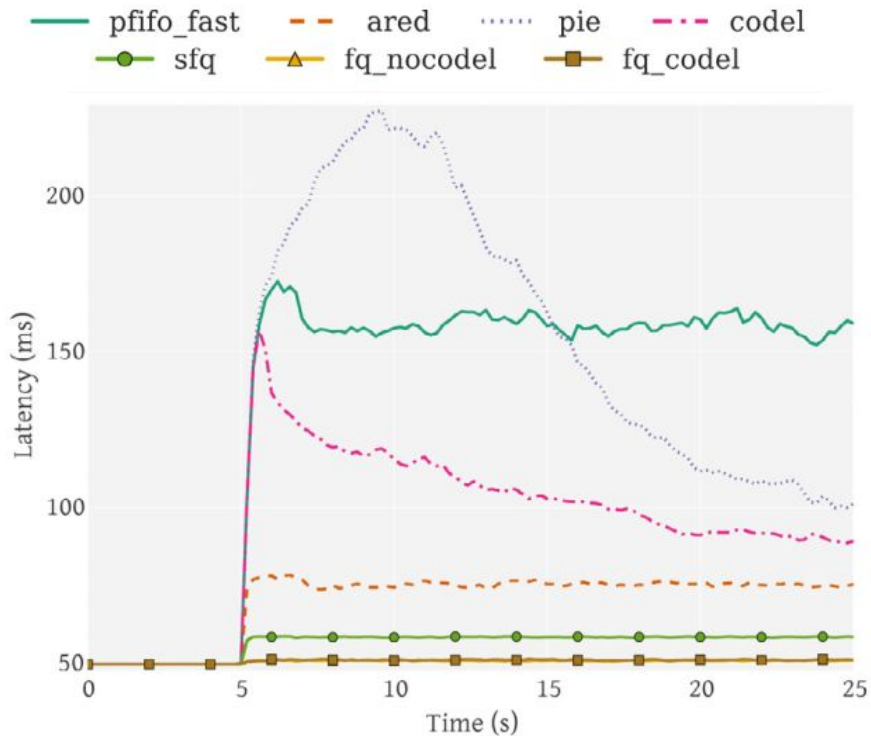
RRUL 10/10



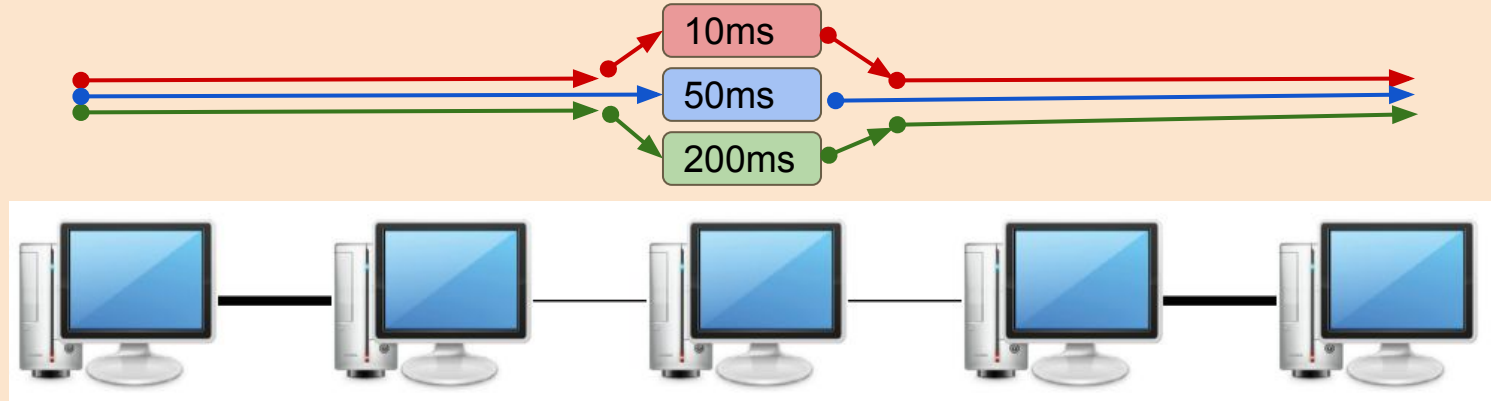
Transient 10/1



Transient 10/10



Fairness Test Implementation



Client:

Iperf -c server
on 3 ports.
Measure per-flow
throughput.

Router1:

```
tc qdisc  
add dev  
eth2 ared
```

Router2:

Use **netem**
for diff
delays.


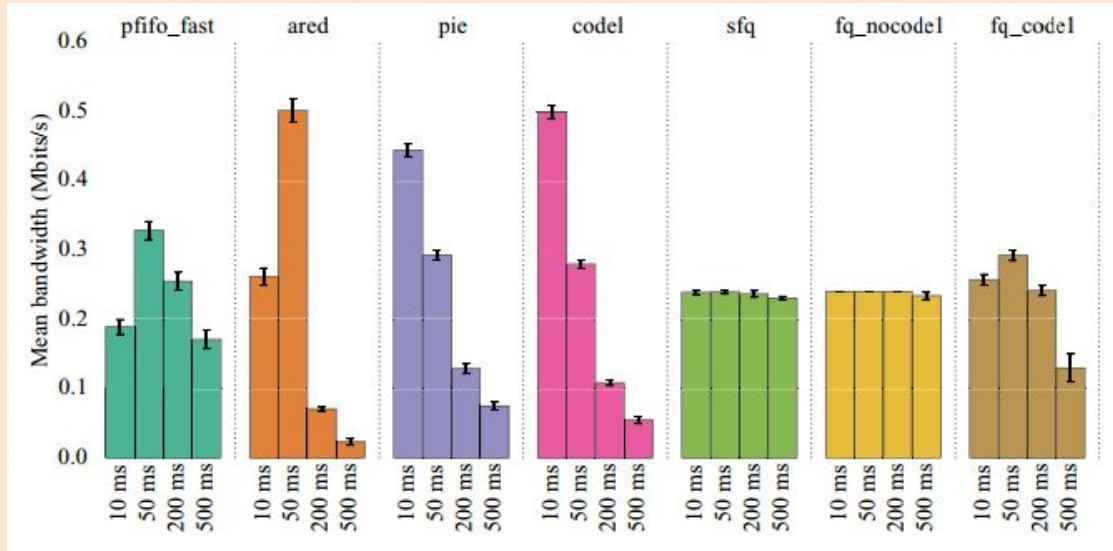
Router3:

```
tc qdisc  
add dev  
eth1 ared
```

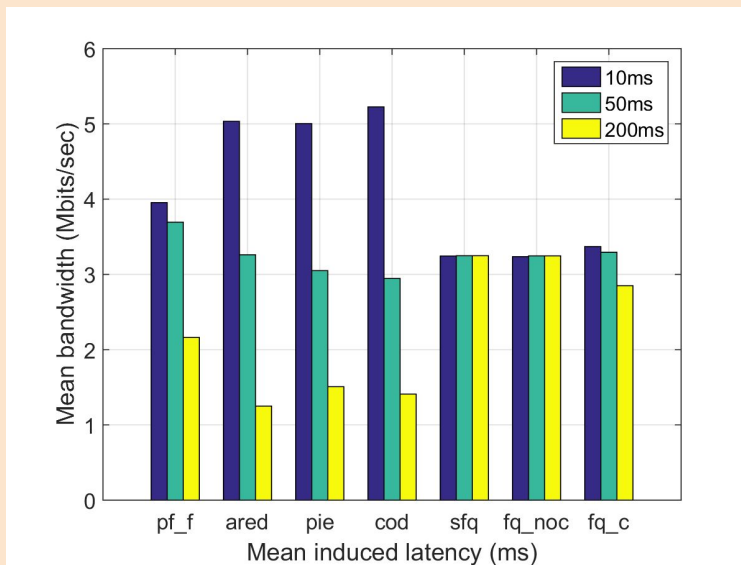
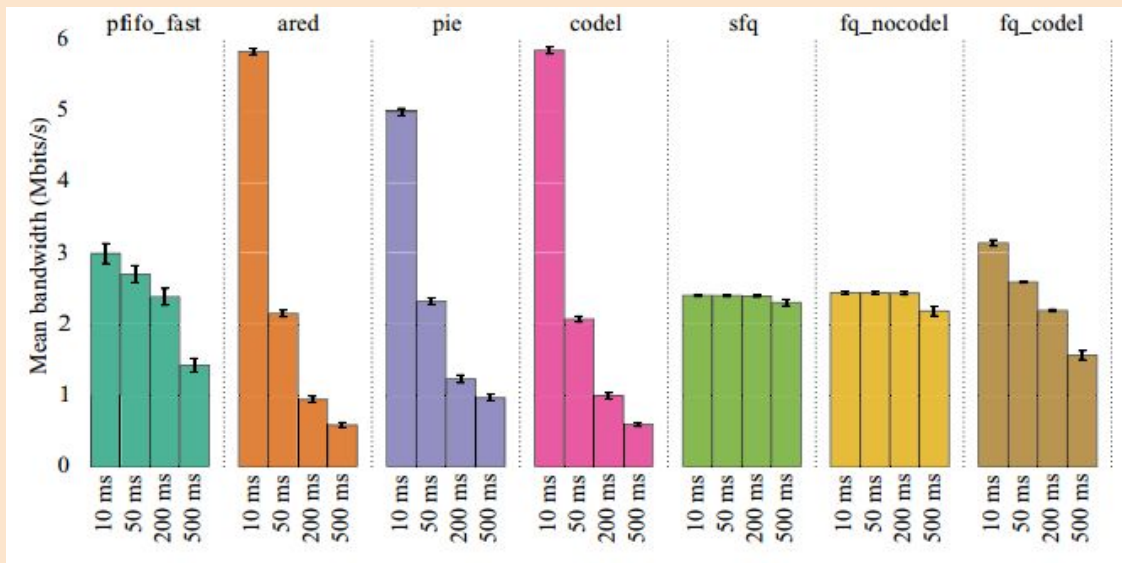
Server:

Iperf -s
on 3 ports

Fairness 10/1



Fairness 10/10


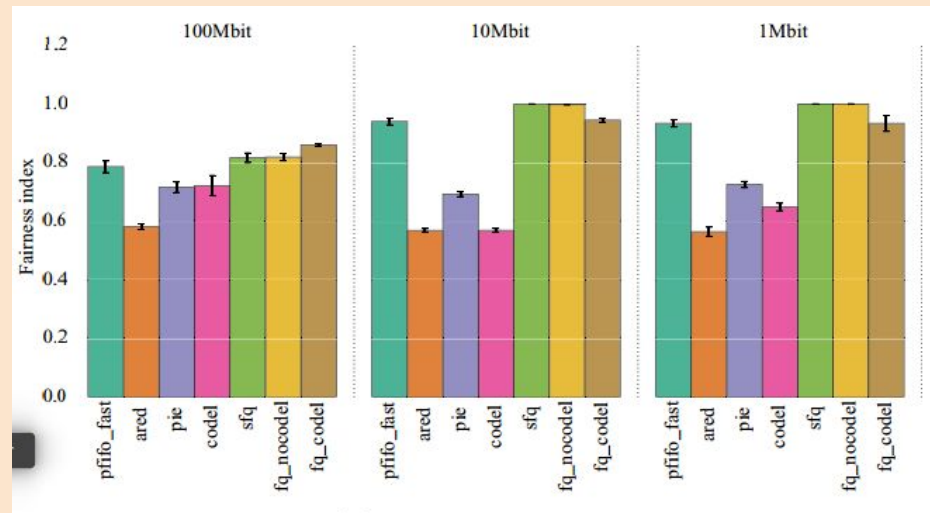


Fairness Index


Fairness Index for 3 flows:

$$\mathcal{L} = \frac{(x_1 + x_2 + x_3)^2}{3(x_1^2 + x_2^2 + x_3^2)}$$


Value between 0 and 1.



WiFi Test



- Throughput Latency graph for **100 Mbps** flows
- Result: same ordering of latency behavior. But, magnitude is different. (Minimum 100ms)
- Reason: Queuing at lower layers (no control in WiFi).



Conclusion

- We compared the performance of **three AQMs** and fairness queueing.
- Good: AQMs perform better than FIFO in terms of **throughput**, and **latency**.
- Bad: AQMs give a **spike in latency** at transient phase and they have a tendency to exacerbate the **unfairness** issue.
- WiFi: AQMs failed to perform better.

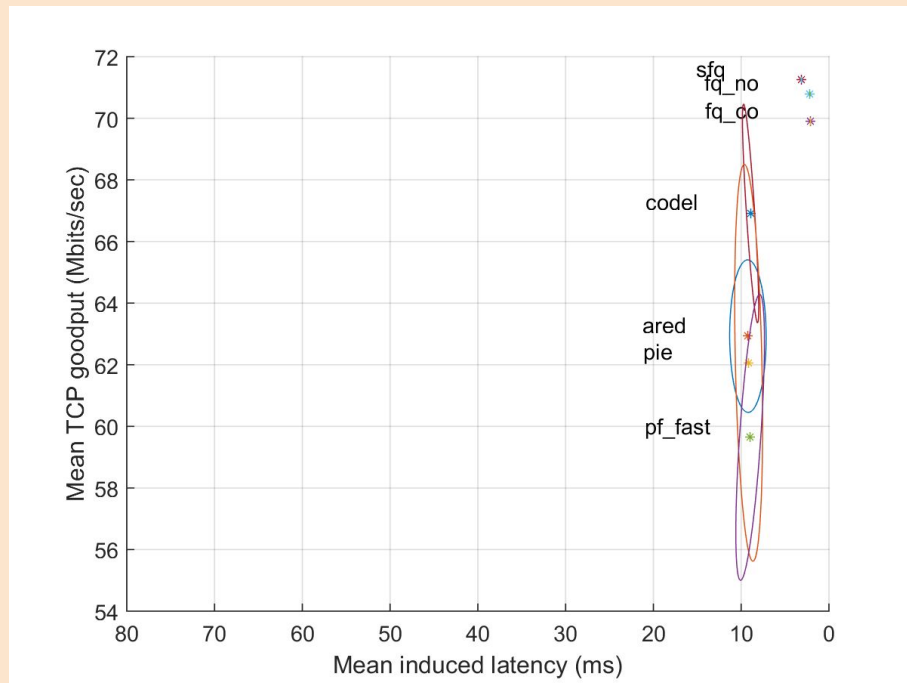

Thank You

Questions!!


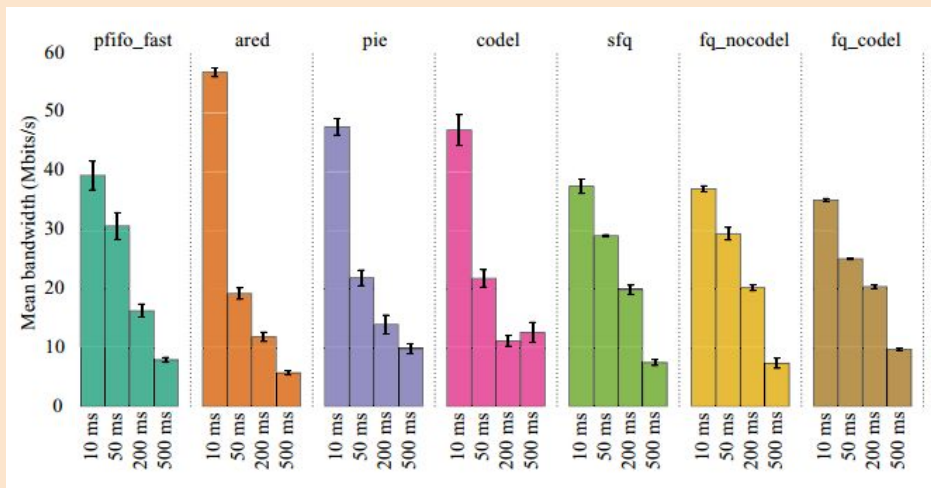
BackUp Slides

100/100 Mbps Flow

RRUL 100/100



Fairness 100/100



Transient 100/100

